

# BioAssist: Biomedical Web-Search Enabled RAG Chatbot

## 1. Project Overview

BioAssist is a domain-specific AI-powered Retrieval-Augmented Generation (RAG) chatbot built to assist users with biomedical, clinical, and healthcare-related queries. It combines knowledge retrieval from internal documents with real-time web search. The chatbot leverages transformer-based sentence embeddings and integrates them into a FAISS-based vector database for efficient similarity search. Its prompt-building strategy incorporates both past conversational history and dual-contextual grounding (local + web), enhancing the relevance and coherence of the generated answers.

---

## 2. Document Ingestion & Vectorization

- **Supported Formats:** PDF, DOCX, TXT, CSV
    - A total of 17 local documents were ingested in the vector index. 5 PDF, 4 CSV, 4 TXT and 4 DOCX
  - **Text Chunking:** Token-based splitter with overlap
  - **Embedding Model:** sentence-transformers/all-MiniLM-L6-v2
  - **Vector DB:** FAISS
- 

## 3. Architecture Overview

The BioAssist architecture follows a modular design enabling efficient ingestion, retrieval, generation, and interaction. It comprises the following key components:

- **Document Ingestion & Vector Indexing**  
Documents in formats like PDF, DOCX, TXT, and CSV are parsed, chunked using a token-based splitter (TokenTextSplitter), and embedded using HuggingFace's sentence-transformers/all-MiniLM-L6-v2. These embeddings are L2-normalized and stored in a FAISS index using IndexFlatIP for cosine similarity-based retrieval. Metadata is attached to each chunk to enable traceability and quality inspection. The pipeline uses LangChain methods to process data.
- **Retrieval & Query Processing**  
Each user query is first evaluated by a domain-specific guardrail to determine if it falls within the biomedical domain. If it does not, the chatbot responds with a fallback message indicating that the query is outside its scope. If the query is biomedical, it is passed to a query rewriter powered by an LLM, which resolves ambiguities (such as pronouns or incomplete references) by incorporating context from the ongoing conversation. The rewritten query is then processed through a conversational retrieval chain, ensuring that responses are both

- context-aware and grounded in relevant biomedical knowledge. The system retrieves the top-5 most relevant document chunks from the FAISS vector store using similarity scores.
- If the best similarity score is below 0.4, it enters web-fallback mode, where only top-5 web search results are considered.
  - Otherwise, it combines both local document chunks and web results for grounding.
  - Web search is executed in every case, ensuring robustness even when local content is available. Retrieval logic is implemented in `rag_pipeline.py`.
  - LangChain is used as a utility layer to format prompts, structure documents, and wrap Gemini model calls.
- **Web Search Integration**

DuckDuckGo web results are fetched using the `ddgs` library and converted into Document objects via LangChain's schema. These are included in the final context for answer generation. In fallback mode (when local similarity is < 0.4), only web search results are used to generate the answer.

- **Prompt Construction & LLM Answer Generation**

The final prompt includes:

- User query
  - Chat history
  - Retrieved local document chunks
  - Web search snippets
  - This is passed to Gemini Flash 2.0 Lite, configured with a temperature of 0.4 and max tokens of 2000. The model generates a grounded, context-aware response that is verified against the retrieved context for factual grounding (Hallucination Detection).
- 

## 4. RAG Pipeline

- **Query Rewriting:** Ensures clarity and self-containment for web search
  - **Biomedical Guardrail:** Binary classifier using Gemini to restrict queries to health-related topics
  - **Dual Retrieval:**
    - From internal vector DB (using similarity search)
    - From DuckDuckGo web search (via `web_search.py`)
  - **Prompt Construction:**
    - Chat history + Local + Web chunks used
    - Final prompt sent to Gemini model for answer generation
  - **Answer Filtering:**
    - Verify grounding of each sentence using self-critique technique. Basically, the generated answer is compared with the provided context.
- 

## 5. Web Search Integration

- **Library Used:** DuckDuckGo (`ddgs`) via LangChain schema

- **Fallback Trigger:** Local similarity score below threshold (e.g., 0.4)
  - **Web Context:** Injected into prompt along with local context
  - **Source Tracking:** All URLs and titles retained to be displayed as reference with the provided answer
- 

## 6. Frontend Interface

- **Framework:** Streamlit
  - **Key Features:**
    - Realtime input/output
    - Exportable chat history
    - End chat option
    - Source traceability (local vs. web)
    - Visual indicators for fallback cases
    - Chat input with auto-scroll
    - Export chat (.txt, .md)
    - Performance metrics (retrieval time, generation time, token count)
    - Past conversation management
    - Web fallback notice and content reference
- 

## 7. Bonus Features Implemented

- **Use of LangChain**
  - **DuckDuckGo-based web search fallback**
  - **Query rewrite using chat context**
  - **Biomedical safety guardrail**
  - **Hallucination Detection**
  - **Streamlit UI with export, metrics, chat persistence**
  - **Logging through “loguru”**
  - **Docker file**
- 

## 8. Design Decisions

- **Gemini Flash 2.0 Lite:** Chosen for context length, speed and output quality.
  - **MiniLM Embeddings:** Small, fast, and high-performing for semantic similarity
  - **FAISSDB:** Configurable backend for flexibility and persistence
  - **DDG Web Search:** Free and fast, avoids dependency on external API keys
-